RELAY

Incident triage for logistics

BUSINESS SUMMARY / FRESH SYNTHETIC TEST / 22 OCTOBER 2026

The decision

Keep routing **human-confirmed**. Text / Sensitive is the local demo default. On 60 new reports it caught 14/15 urgent cases, but missed a critical fire-alarm blocker and flagged 50/60 for review. This supports a work sample, not automatic operational routing.

Problem and working prototype

Relay proposes a receiving team and urgency for logistics incidents, flags uncertainty and records coordinator corrections. Compact queues lead from report to confirmation and Done; model controls and evidence stay in Evaluation. Shared checks compare stated temperature limits and readings and flag active exposures for review. No real dispatch is sent.

Evidence

Two approaches: rules and trained TF-IDF centroids, with shared cold-chain protection. There are 300 development reports (240 training, 48 families), 60 original test cases and 60 new scenarios. Before the new set was authored, the classifier and Text/Sensitive choice were frozen. Another fresh-context agent saw only new text and the unchanged guide. Labels were frozen before scoring; no tuning followed.

Approach / policy	Macro F1	Urgent recall	Review	Cost*	p95 ms
Rules / Standard	0.571	20%	39/60	1281	0.154
Rules / Sensitive	0.571	60%	51/60	687	0.160
Text / Standard	0.600	80%	34/60	384	0.066
Text / Sensitive	0.600	93%	50/60	164	0.059

*Relative cost with miss penalty 100 for every row. Category: 56 known labels; urgency: 54 known, including 15 urgent. Unknowns remain unknown. Review uses all 60. Team accuracy is 61%; raw scores include reviewed cases. p95 excludes HTTP, startup and humans.

Changed requirement and workload

At a common missed-urgent penalty of 100, Text / Standard to Sensitive changes misses 3 to 1, reviews 34 to 50, and cost 384 to 164. Text / Sensitive has 7 false urgent alerts, 0 wrong eligible teams and 0 bypassed required reviews. Only ten reports are eligible; this does not establish automatic-routing reliability. RLY-910052 still gets routine customs review despite a disabled building fire alarm.

Economics, limits and next step

Both classifiers incur **USD 0 in inference API calls**; compute, development/review services and labor are excluded. Relative assumptions: review 1, wrong eligible route or bypassed required review 6 once, false urgency 2, missed priority 30 or 100 even during review. These are assumptions, not savings. Synthetic case selection, author knowledge of the classifier, shared model-family bias and guide ambiguity limit the fresh evidence.

Next: clarify critical safety consequences with an operator, obtain representative reports and domain review, measure review capacity, and use another untouched set for any tuned version. No field trial has occurred.

Evidence: evidence/fresh-20261022/fresh-test.json; FRESH-VALIDATION.md; EVALUATION.md. Manual prototype testing and feedback: dansyar (author-reported). Software, automated checks and evaluation: OpenAI Codex. Reference labels: separate Codex agents. Full roles in ATTRIBUTION.md.