RELAY

Incident triage for logistics

BUSINESS SUMMARY / FRESH 2 SYNTHETIC TEST / 22 OCTOBER 2026

The decision

Keep routing **human-confirmed**. Text / Sensitive is the local demo default. On the second new set, it gets 49/55 clear team labels correct (89%) and catches all 13 labeled urgent cases, while reviewing 43/60. One wrong team outside model abstention remains.

Problem and working prototype

Relay proposes teams and urgency, flags uncertainty and records coordinator corrections. Compact queues lead from report to confirmation and Done; controls and evidence stay in Evaluation. The revision checks primary issues, cargo custody and active consequences. The missed fire-alarm controller now receives Safety/Urgent review in all four configurations. No real dispatch is sent.

Evidence

Rules and TF-IDF centroids share the new evidence checks; training stays at 240 reports. There are 480 fictional reports, including three independently agent-labeled sets of 60. The revised classifier and Text/Sensitive choice were frozen before Fresh 2 was authored. A new fresh-context agent saw only new text and the unchanged guide. Labels were frozen before scoring; no tuning followed.

Approach / policy	Macro F1	Urgent recall	Review	Cost*	p95 ms
Rules / Standard	0.768	69%	36/60	454	0.205
Rules / Sensitive	0.768	85%	54/60	272	0.195
Text / Standard	0.904	100%	30/60	54	0.109
Text / Sensitive	0.904	100%	43/60	59	0.110

*Relative cost with miss penalty 100 for every row. Category: 55 known labels; urgency: 52 known, including 13 urgent. Unknowns remain unknown. Review uses all 60. Team accuracy is 89%; raw scores include reviewed cases. p95 excludes HTTP, startup and humans.

Changed requirement and workload

On the same Fresh 2 reports, preserved v3 to v4 improves teams 46/55 to 49/55 (83.6% to 89.1%), F1 0.839 to 0.904 and false urgent alerts 7 to 5; reviews rise 41 to 43. Both catch all 13 urgent cases. This paired comparison was added after primary scoring, without tuning. Three remaining team disagreements concern operational versus Safety ownership of urgent blockers.

Economics, limits and next step

At miss penalty 100, Text/Standard to Sensitive costs 54 to 59 and reviews 30 to 43, catching 8 to 10 required reviews. Sensitive is not cheapest on this set; it remains the preregistered conservative default. Review costs 1, wrong eligible team or bypassed review 6 once, false urgency 2, and every urgent miss 30 or 100 even in ordinary review. **USD 0 inference API calls** excludes compute, review/development services and labor; relative costs are assumptions, not savings.

Synthetic selection, author knowledge, shared model-family bias and guide ambiguity limit the evidence. Next: operator guidance on critical-system team ownership, representative reports and measured review capacity. Any further tuned version needs another untouched set. No field trial has occurred.

Evidence: evidence/fresh2-20261022/fresh-test.json; FRESH2-VALIDATION.md; EVALUATION.md. Manual prototype testing and feedback: dansyar (author-reported). Software, automated checks and evaluation: OpenAI Codex. Reference labels: separate Codex agents. Full roles in ATTRIBUTION.md.